

# Graded Project Proposal

- Monday, September 26
- PDF document (1-2 pages) to be submitted on blackboard
  - Team members
  - Motivation & Problem Statement
  - Approach
  - Deliverables & Timeline
- Things to keep in mind
  - Define a problem you can reasonably tackle in one semester
    - Of a backup plan if you want to try something more ambitious
  - Data collection methodologies (if you are using data)
- In-class presentation by members
  - PowerPoint, other props allowed

# Using Markov Decision Processes for Policy Enforcement

September 16, 2016

# Automated enforcement of privacy policies

- Healthcare industry has had numerous high profile privacy incidents
- Government is increasingly cracking down on such breaches
- Privacy is an *ongoing process*
  - Systematically identify employees who are engaging in risky behaviors indicative of snooping, identity theft, etc
  - Must be done *regularly* on *all employees* with access to protected information
- Manual audit log reviews are tedious and expensive

# Verifying an organization obeys privacy policies

- Privacy policies restrict the purpose for which some protected information can be used for
- Types of restrictions
  - Prohibitive rules
    - An organization does *not* use certain information *for* a purpose
    - Example: Yahoo!'s privacy policy
  - Exclusivity rules
    - An organization users certain information *only for* a given list of purposes
    - Example: HIPAA

# Enforcing privacy policies

- Automating enforcement is desirable
- Problem:
  - Tools for enforcement assume that actions they audit are **accurately** labeled with the purpose
  - Tools do not analyze the meaning of *purpose*
- Solution:
  - Need a way to determine if an action could be for a purpose

# Purpose in policy enforcement


- Planning is central to purpose
  - Cognitive psychology: Purpose is the central determinant of behavior
- Formalizing the relationship between planning and purpose
  - Use MDP to construct all all possible behaviors a privacy policy allows
  - Auditing an organization:
    - MDP model of allowed behaviors VS
    - All behaviors the organization being audited that could result in the given audit log

# Planning for a purpose

- Model the environment the organization being audited operates as an MDP
- Why MDP?
  - Good for planning with probabilistic systems
  - Reward function quantifiers the degree of progress towards a purpose upon taking an action from a state
- If the organization being audited is action only to achieve a particular purpose:
  - Actions correspond to an optimal plan for that MDP
  - Actions are for that purpose

# MDP: Formally


- $MDP$  is a tuple  $\langle S, A, T, R, \gamma \rangle$ 
  - $S$  is a finite state space
  - $A$  is a finite action set
  - $T: S \times A \rightarrow D(S)$  is a transition function that given a state and an action, maps to a distribution over states  $D(S)$
  - $R: S \times A \rightarrow \mathbb{R}$  is a reward function
  - $\gamma$ : a discount factor such that  $0 < \gamma < 1$

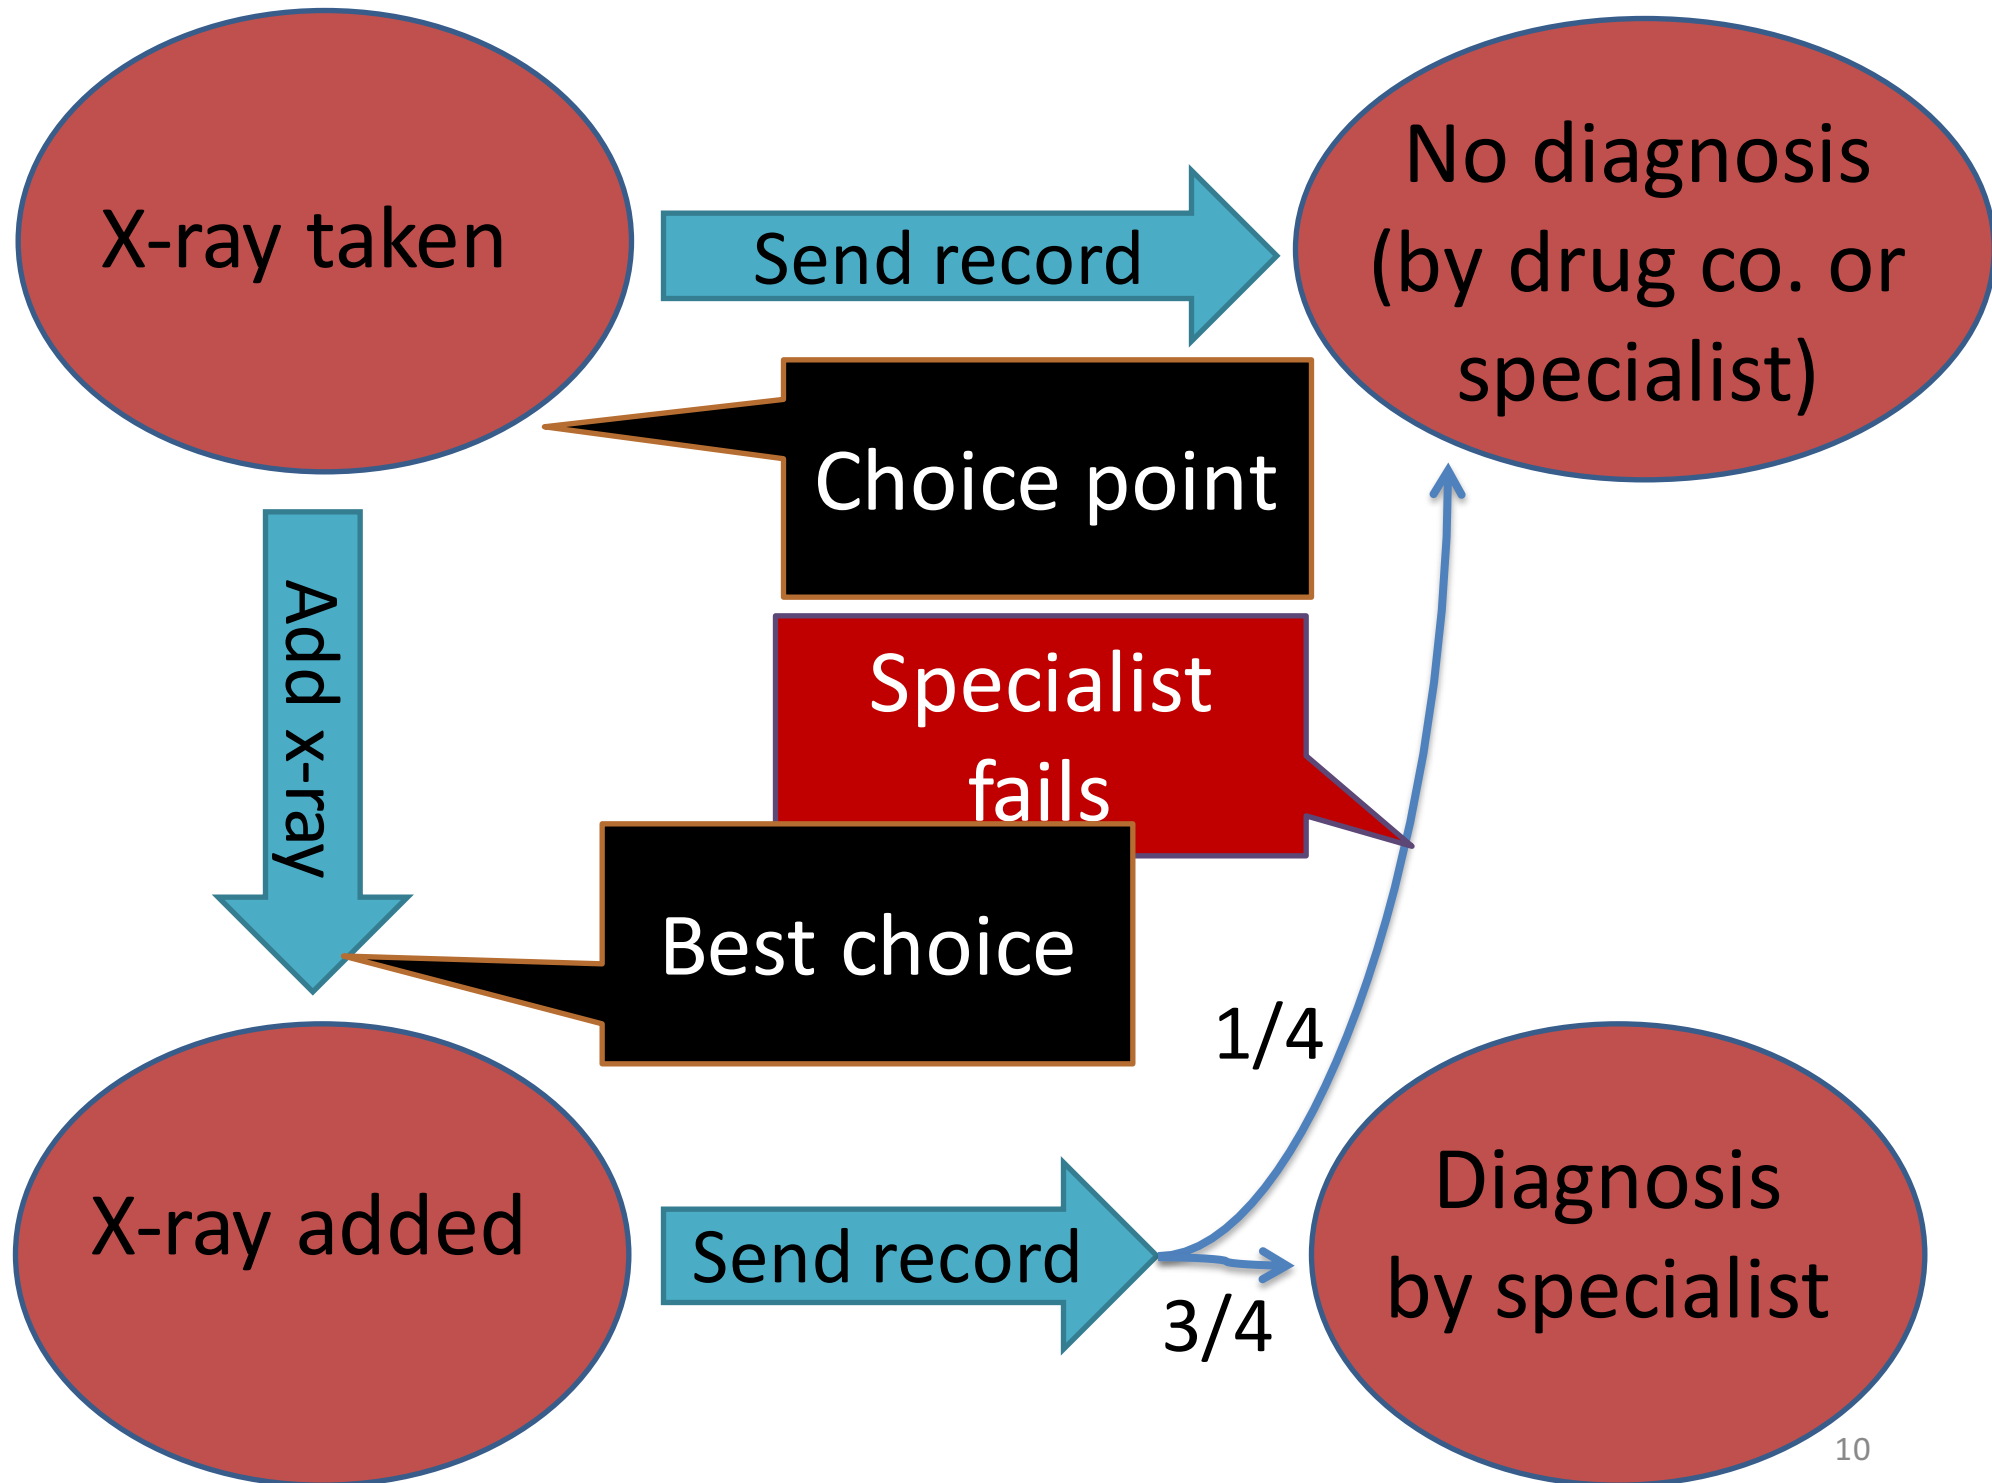




# Example process modeled by MDP

- Actions: Left, Right, Up, Down
- Entering top-right corner gives reward of +1 and then takes you to a random state.
- Finding good strategy that maximizes reward.





# Goal of an agent using MDP to plan its actions

- $MDP = \langle S, A, T, R, \gamma \rangle$
- ***Maximize its expected total discounted reward  $\mathcal{E}[\sum_{i=0}^{\infty} \gamma^i r_i]$*** 
  - $i$  is a natural number ranging over time modeled as discrete steps
  - $r_i$  is the reward at time  $i$
  - Expectation taken over probabilistic transitions
  - Discount factor  $\gamma$  accounts for preference of people to receive rewards sooner, not later

# What is a good strategy?

- What we want a strategy to optimize:
  - Expected reward / time step
  - Expected reward in the first  $t$  steps

# Maximizing the discounted reward

$$Q(s, a) = r_0 + \gamma r_1 + \gamma^2 r_2 + \gamma^3 r_3 + \dots$$

$$V(s) = \max_a Q(s, a)$$

$$V(s) = \max_a [R(s, a) + \gamma \sum_{s'} P(s') V(s')]$$

# Solving for the maximum discounted award

- Given:
  - Transition function
  - Reward function
- Possible Solutions
  - Dynamic programming
    - Start with guesses  $V_0(s)$  for all states  $s$
    - Update using 
$$V_i(s) = \max_a \left[ E[R(s, a)] + \gamma \sum_{s'} \text{Pr}(s') V_{i-1}(s') \right]$$
  - Linear programming
    - Replace “max” with “ $\geq$ ” and minimize  $\sum_s V(s)$  subject to the constraints

# Infinite MDP

There exists an optimal stationary policy, with certain conditions on rewards.

# Was that an infinite MDP?

- How can you make it one?
  - Add a stop action that loops at the end



# Audit Algorithm

```
AUDIT( $m = \langle \mathcal{S}, \mathcal{A}, t, r, \gamma \rangle$ ,  $b = [s_1, a_1, \dots, s_n, a_n]$ ):  
01  if (IMPOSSIBLE( $m, b$ ))  
02      return true      // behavior impossible for NMDP  
03   $V_m^* := \text{SOLVEMDP}(m)$   
04  for ( $i := 1$ ;  $i \leq n$ ;  $i++$ ):  
05      if ( $Q^*(V_m^*, s_i, a_i) < V_m^*(s_i)$ ):  
06          return true      // action suboptimal  
07      if ( $Q^*(V_m^*, s_i, a_i) \leq 0$  and  $a_i \neq \text{Stop}$ ):  
08          return true      // action redundant  
09  return false
```

Figure 2. The algorithm AUDIT

Questions on the homework?